

Linux /dev/random

A New Approach

Stephan Müller
<smueller@chronox.de>

Agenda

- LRNG Goals
- LRNG Design
- Entropy Sources
- Initial Seeding Strategies

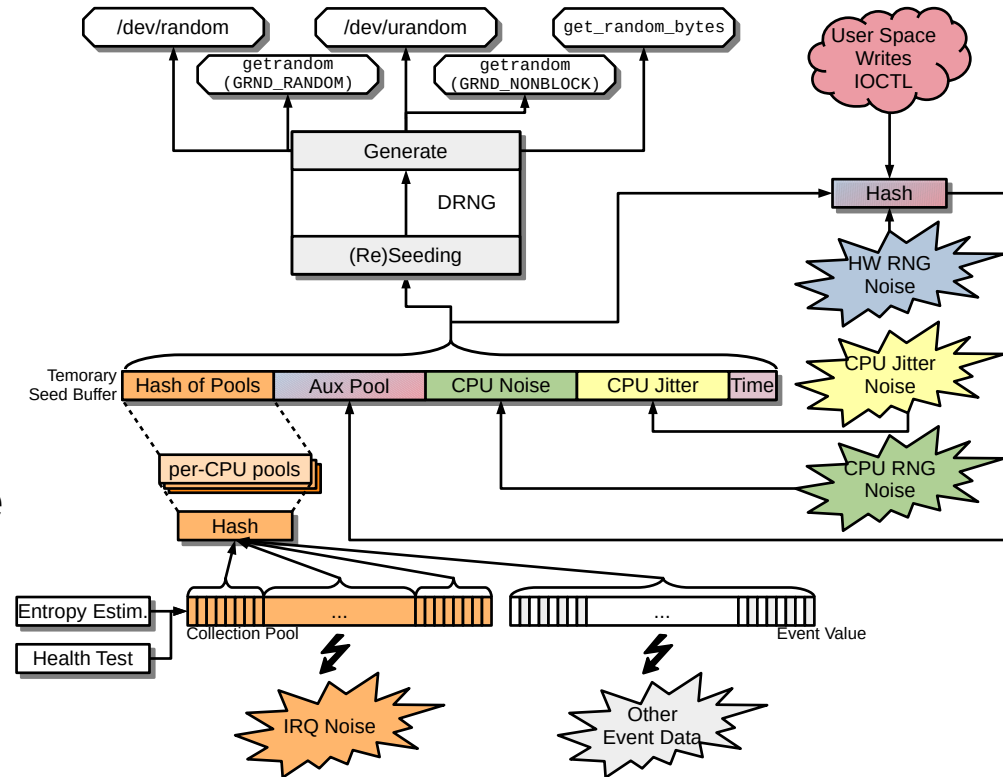
LRNG Goals

- Sole use of cryptography for data processing
- High-Performance lockless IRQ handler
- Test interfaces for all LRNG processing steps
- Power-up and runtime tests
- API and ABI compliant drop-in replacement of existing /dev/random
- Flexible configuration supporting wide range of use cases
- Runtime selection of cryptographic implementations
- Clean architecture
- Standards compliance: SP800-90A/B/C, AIS 20/31

```
-- Linux Random Number Generator
Specific DRNG seeding strategies --->
Entropy Source Configuration --->
[*] Support DRNG runtime switching --->
[*] LRNG testing interfaces --->
[*] Enable power-on and on-demand self-tests
[ ] Panic the kernel upon self-test failure
```

LRNG Design

- 4 Entropy Sources
 - 3 external
 - 1 internal
 - All ES treated equally
 - No domination by any ES – seeding triggered by boot process or DRNG
- All ES can be selectively disabled at runtime
- ES data fed into DRNG
- DRNG accessible with APIs



DRNG Output APIs

- Blocking APIs – deliver data only after fully initialized and fully seeded:
 - `/dev/random`
 - `getrandom()` system call
 - `get_random_bytes_full` in-kernel API after being triggered with `add_random_ready_callback`
- All other APIs deliver data without blocking until complete initialization
 - No guarantee of LRNG being fully initialized / seeded


DRNG Seeding

- Temporary seed buffer: concatenation of output from all ES
- Seeding during boot: when 32/128/256 bits of entropy are available
- Seeding at runtime:
 - After 2^{20} generate requests or 10 minutes
 - After forced reseed by user space
 - At least 128 bits (SP800-90C mode: LRNG security strength) of total entropy must be available
 - 256 bits of entropy requested from each ES – ES may deliver less
 - Seed operation occurs when DRNG is requested to produce random bits

```
lrng_pool: obtained 192 bits by collecting 256 bits of entropy from aux pool, 0 bits of entropy remaining
lrng_sw_noise: 256 interrupts used from entropy pool of CPU 4, 256 interrupts remain unused
lrng_sw_noise: 0 interrupts used from entropy pool of CPU 5, 512 interrupts remain unused
lrng_sw_noise: 0 interrupts used from entropy pool of CPU 15, 174 interrupts remain unused
lrng_sw_noise: obtained 192 bits by collecting 256 bits of entropy from entropy pool noise source
lrng_archrandom: obtained 384 bits of entropy from CPU RNG noise source
lrng_jent: obtained 24 bits of entropy from Jitter RNG noise source
lrng_drng: seeding regular DRNG with 212 bytes
lrng_drng: regular DRNG stats since last seeding: 1397 secs; generate calls: 1
```

DRNG Management

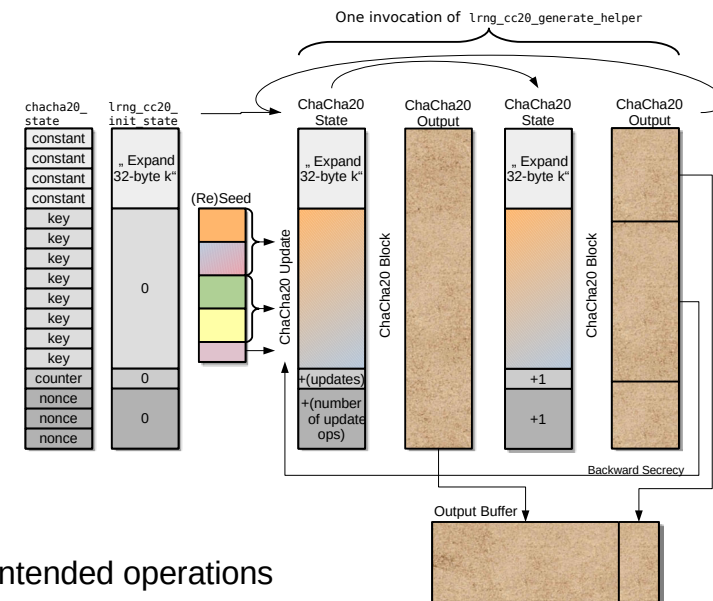
- One DRNG per NUMA node
- Hash contexts NUMA-node local
- Each DRNG initializes from entropy sources
- Sequential initialization of DRNG – first is Node 0
- If DRNG on one NUMA node is not yet fully seeded → use of DRNG(Node 0)
- Each DRNG instance managed independently
- To prevent reseed storm – reseed threshold different for each node
 - Node 0: 600 seconds
 - Node 1: 700 seconds
 - ...
- NUMA support code only compiled if CONFIG_NUMA → only one DRNG present



Data Processing Primitives

- Sole use of cryptographic mechanisms for data compression
- Cryptographic primitives Boot-Time / Runtime switchable
 - Switching support is compile-time option
 - DRNG, Conditioning hash
 - Built-in: ChaCha20 DRNG / SHA-256
 - Available:
 - SP800-90A DRBG (CTR/Hash/HMAC) using accelerated AES / SHA primitive, accelerated SHA-512 conditioning hash
 - Hardware DRNG may be used (e.g. CPACF)
 - Well-defined API to allow other cryptographic primitive implementations
- Complete cryptographic primitive testing available:
 - Full ACVP test harness available: <https://github.com/smuellerDD/acvpparser>
 - ChaCha20 DRNG userspace implementation: https://github.com/smuellerDD/chacha20_drng
- Other data processing primitives:
 - Concatenation of data
 - Truncation of message digest to heuristic entropy value
- Entropy behavior of all data processing primitives based on fully understood and uncontended operations


- Support DRNG runtime switching
<M> SP800-90A support for the LRNG
<M> Kernel Crypto API support for the LRNG



External Entropy Sources


- Use without additional conditioning – fast source
 - Jitter RNG
 - CPU (e.g. Intel RDSEED, POWER DARN, ARM SMC Calling Convention or RNDR register)
 - Data immediately available when LRNG requests it
- Additional conditioning – slow source
 - RNGDs
 - In-kernel hardware RNG drivers
 - All received data added to “auxiliary pool” with hash update operation
 - Data “trickles in” over time
- Every entropy source has individual entropy estimate
 - Taken at face value – each ES requires its own entropy assessment

```
*** Jitter RNG Entropy Source ***  
[*] Enable Jitter RNG as LRNG Seed Source  
(16) Jitter RNG Entropy Source Entropy Rate  
*** CPU Entropy Source ***  
(8) CPU Entropy Source Entropy Rate
```



Internal Entropy Source

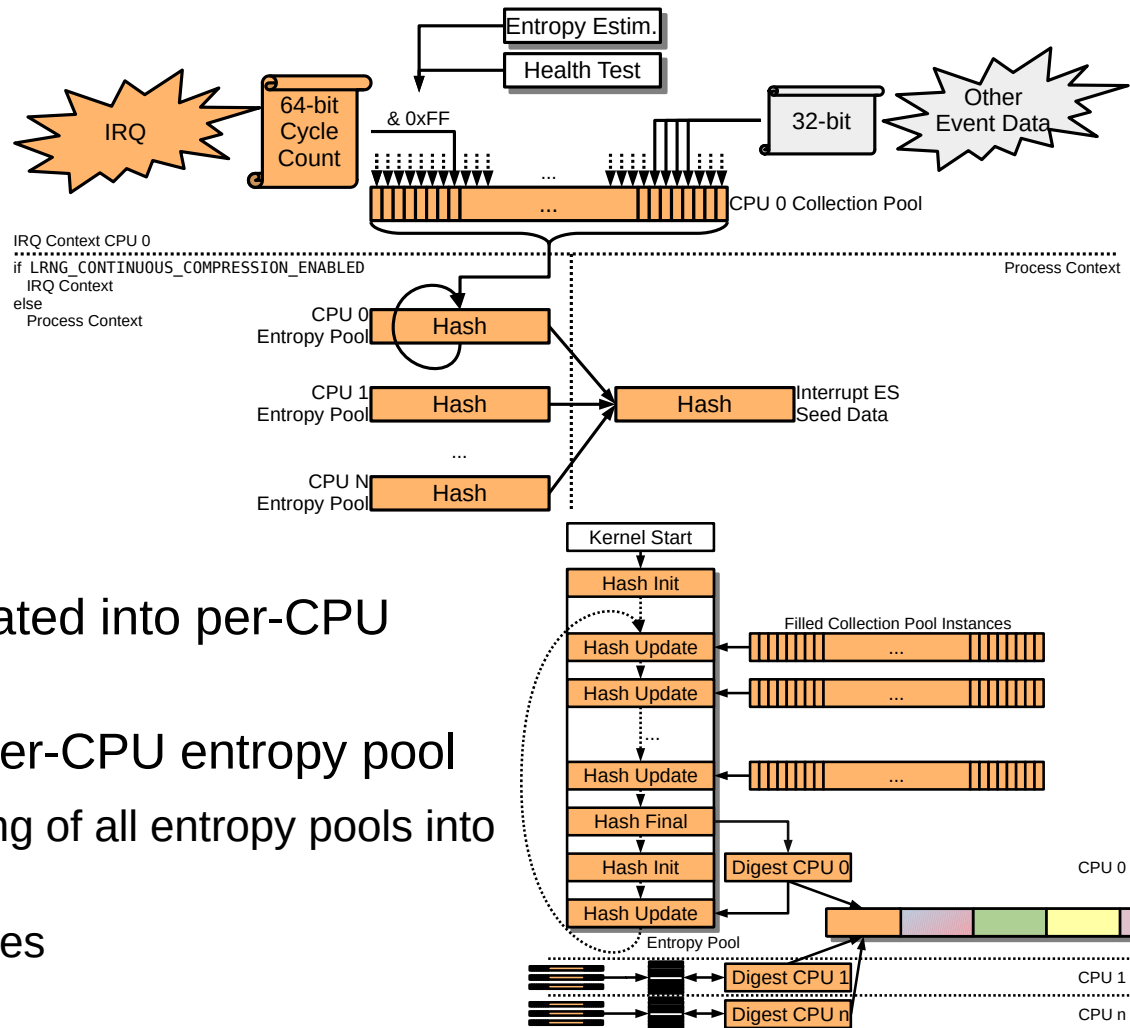
- Interrupt timing
 - All interrupts are treated as one entropy source
- Data collection executed in IRQ context
- Data compression executed partially in IRQ and process context
- Data compression is a hash update operation
- High performance: up to twice as fast as legacy /dev/random in IRQ context with LRNG_CONTINUOUS_COMPRESSION enabled
 - Even faster without continuous compression



```
*** Interrupt Entropy Source ***
Continuous entropy compression boot time setting (Enable cont
[*] Runtime-switchable continuous entropy compression
LRNG Entropy Collection Pool Size (1024 interrupt events (def
[*] Enable interrupt entropy source online health tests
(256) Interrupt Entropy Source Entropy Rate
```

Internal ES Data Processing

- 8 LSB of time stamp concatenated into per-CPU collection pool
 - Entropy estimate
 - Health test
- 32 bits of other event data concatenated into per-CPU collection pool
- When array full → conditioned into per-CPU entropy pool
 - When entropy is required → conditioning of all entropy pools into one message digest
 - Addition of all per-CPU entropy estimates



Internal ES Testing Interfaces

- Testing code is compile time option
- Access via DebugFS
- Testing supports data collection at boot time and runtime:
 - Raw unprocessed entropy time stamps for IRQ
 - Raw Jiffies for IRQ
 - IRQ value
 - IRQ flags value
 - _RET_IP_ per IRQ
 - Performance data for LRNG's IRQ handler
- Hash testing interface for built-in SHA-256
- Full SP800-90B assessment documentation
- Raw entropy collection and analysis tools provided

Test System	Entropy of 1,000,000 Traces	Sufficient Entropy
ARMv7 rev 5	1.9344	Y
ARMv7 rev 5 (Freescale i.MX53) ^[22]	7.07088	Y
ARMv7 rev 5 (Freescale i.MX6 Ultralite) ^[23]	6.638399	Y
ARM 64 bit AppliedMicro X-Gene Mustang Board	5.599128	Y
Intel Atom Z530 – using GUI	3.38584	Y
Intel i7 7500U Skylake - 64-bit KVM environment	3.452064	Y
Intel i7 8565U Whiskey Lake – 64-bit KVM environment	7.400136	Y
Intel i7 8565U Whiskey Lake – 32-bit KVM environment	7.405704	Y
Intel i7 8565U Whiskey Lake	6.871	Y
Intel Xeon Gold 6234	4.434168	Y
IBM POWER 8 LE 8286-42A	6.830712	Y
IBM POWER 7 BE 8202-E4C	4.233912	Y
IBM System Z z13 (machine 2964)	4.366368	Y
IBM System Z z15 (machine 8561)	5.691832	Y
MIPS Atheros AR7241 rev 1 ^[24]	7.157064	Y
MIPS Lantiq 34Kc V5.6 ^[25]	7.032740	Y
Qualcomm IPQ4019 ARMv7 ^[26]	6.638405	Y
SiFive HiFive Unmatched RISC-V U74	2.387470	Y

-LRNG testing interfaces

```
[*] Enable entropy test interface to hires timer noise source
[*] Enable entropy test interface to Jiffies noise source
[*] Enable entropy test interface to IRQ number noise source
[*] Enable entropy test interface to IRQ flags noise source
[*] Enable entropy test interface to RETIP value noise source
[*] Enable entropy test interface to IRQ register value noise source
[*] Enable test interface to LRNG raw entropy storage array
[*] Enable LRNG interrupt performance monitor
[*] Enable LRNG ACVT Hash interface
[*] Enable runtime configuration of entropy sources
```

Internal ES Health Test

- Health test compile-time configurable
- Power-Up self tests
 - All cryptographic mechanisms
 - Time stamp management
- APT / RCT
- Time-Stamp Pattern detection: 1st/2nd/3rd discrete derivative of time $\neq 0$
- Blocking interface: Wait until APT power-up testing complete
- Provides SP800-90B compliance of internal ES

CONFIG_LRNG_SELFTEST:

The power-on self-tests are executed during boot time covering the ChaCha20 DRNG, the hash operation used for processing the entropy pools and the auxiliary pool, and the time stamp management of the LRNG.

The on-demand self-tests are triggered by writing any value into the SysFS file `selftest_status`. At the same time, when reading this file, the test status is returned. A zero indicates that all tests were executed successfully.

CONFIG_LRNG_HEALTH_TESTS:

The online health tests validate the noise source at runtime for fatal errors. These tests include SP800-90B compliant tests which are invoked if the system is booted with `fips=1`. In case of fatal errors during active SP800-90B tests, the issue is logged and the noise data is discarded. These tests are required for full compliance with SP800-90B.

Initial Seeding Strategy I

Default Operation

- DRNG is initially seeded with at least 32 bits of entropy
- DRNG is minimally seeded with at least 128 bits of entropy
- DRNG is fully seeded with 256 bits of entropy
- Blocking interfaces released after DRNG is fully seeded
- Default applied:
 - Either no specific seeding strategy compiled
 - Or specific seeding strategy is not enabled at boottime

```
[ ] Oversample entropy sources  
[ ] AIS 20/31 NTG.1 seeding strategy
```

Initial Seeding Strategy II

Entropy Source Oversampling

- Initial / minimal seeding steps apply – fully seeded step changed
- Compile time option
 - Function only enabled in FIPS mode
 - Function only enabled if message digest of conditioner ≥ 384 bits
- Final conditioning: $s + 64$ bit
- Initial DRNG seeding: every entropy source requested for $s + 128$ bits
 - Every ES alone could provide all required entropy
- All ES data concatenated into seed buffer
- Runtime debug mode: display of all processing steps
- SP800-90C compliance:
 - SP800-90A DRBG with 256 bit strength / SHA-512 vetted conditioning component
 - Complies with RBG2(NP) per default
 - Can be configured to provide RBG2(P)
- Can be used in parallel with seeding strategy III

`CONFIG_LRNG_OVERSAMPLE_ENTROPY_SOURCES:`

When enabling this option, the entropy sources are over-sampled with the following approach: First, the entropy sources are requested to provide 64 bits more entropy than the size of the entropy buffer. For example, if the entropy buffer is 256 bits, 320 bits of entropy is requested to fill that buffer.

Second, the seed operation of the deterministic RNG requests 128 bits more data from each entropy source than the security strength of the DRNG during initialization. A prerequisite for this operation is that the digest size of the used hash must be at least equally large to generate that buffer. If the prerequisite is not met, this oversampling is not applied.

This strategy is intended to offset the asymptotic entropy increase to reach full entropy in a buffer.

The strategy is consistent with the requirements in NIST SP800-90C.

General Testing

- Automated regression test suite covering the different options of LRNG
- Locking torture test of loading/unloading DRNG extensions under full load
- Applied kernel framework tests:
 - KASAN
 - UBSAN
 - Lockdep
 - Memory leak detector
- Use of LRNG without kernel crypto API
- Performance tests of DRNG
- Syscall validation testing
- Test of LRNG behavior in atomic contexts

```
Executing test with kernel command line fips=1 lrng_jent.jitterrng=256 lrng_arch
random.archrandom=256
[PASSED] Jitter RNG: Jitter RNG working on system
[PASSED] Jitter RNG: used for seeding
Executing test with kernel command line lrng_jent.jitterrng=256 lrng_archrandom.
archrandom=256
[PASSED] Jitter RNG: Jitter RNG working on system
[PASSED] Jitter RNG: used for seeding
Executing test with kernel command line lrng_pool.ntgl=1 lrng_jent.jitterrng=256
lrng_archrandom.archrandom=256
[PASSED] Jitter RNG: Jitter RNG working on system
[PASSED] Jitter RNG: used for seeding
Executing test with kernel command line fips=1 lrng_pool.ntgl=1 lrng_jent.jitter
rng=256 lrng_archrandom.archrandom=256
[PASSED] Jitter RNG: Jitter RNG working on system
[PASSED] Jitter RNG: used for seeding
[PASSED] no failures
===== Testing ended Do 10. Jun 11:30:27 CEST 2021 =====
===== Testing started Do 10. Jun 11:30:27 CEST 2021 =====
Executing test with kernel command line fips=1 lrng_jent.jitterrng=256 lrng_arch
random.archrandom=256
[PASSED] Atomic: LRNG executing in atomic contexts
Executing test with kernel command line lrng_jent.jitterrng=256 lrng_archrandom.
archrandom=256
[PASSED] Atomic: LRNG executing in atomic contexts
Executing test with kernel command line lrng_pool.ntgl=1 lrng_jent.jitterrng=256
lrng_archrandom.archrandom=256
[PASSED] Atomic: LRNG executing in atomic contexts
Executing test with kernel command line fips=1 lrng_pool.ntgl=1 lrng_jent.jitter
rng=256 lrng_archrandom.archrandom=256
[PASSED] Atomic: LRNG executing in atomic contexts
[PASSED] no failures
===== Testing ended Do 10. Jun 11:30:44 CEST 2021 =====
[PASSED] ALL TESTS PASSED
[PASSED] no failures
===== Testing ended Do 10. Jun 11:30:44 CEST 2021 =====
```

DRNG Type	Cipher	Cipher Impl.	Read Size	Performance
HMAC DRBG	SHA-512	C	64 bytes	13.8 MB/s
HMAC DRBG	SHA-512	AVX2	16 bytes	4.7 MB/s
HMAC DRBG	SHA-512	AVX2	32 bytes	11.6 MB/s
HMAC DRBG	SHA-512	AVX2	64 bytes	23.3 MB/s
HMAC DRBG	SHA-512	AVX2	128 bytes	38.3 MB/s
HMAC DRBG	SHA-512	AVX2	4096 bytes	92.1 MB/s
Hash DRBG	SHA-512	C	64 bytes	27.9 MB/s
Hash DRBG	SHA-512	AVX2	16 bytes	13.1 MB/s
Hash DRBG	SHA-512	AVX2	32 bytes	25.9 MB/s
Hash DRBG	SHA-512	AVX2	64 bytes	51.1 MB/s
Hash DRBG	SHA-512	AVX2	128 bytes	83.3 MB/s
Hash DRBG	SHA-512	AVX2	4096 bytes	217.8 MB/s
CTR DRBG	AES-256	C	16 bytes	15.4 MB/s
CTR DRBG	AES-256	AES-NI	16 bytes	24.4 MB/s
CTR DRBG	AES-256	AES-NI	32 bytes	49.3 MB/s
CTR DRBG	AES-256	AES-NI	64 bytes	96.2 MB/s
CTR DRBG	AES-256	AES-NI	128 bytes	177.1 MB/s
CTR DRBG	AES-256	AES-NI	4096 bytes	1.247 GB/s
ChaCha20	ChaCha20	C	16 bytes	42.0 MB/s
ChaCha20	ChaCha20	C	32 bytes	84.5 MB/s
ChaCha20	ChaCha20	C	64 bytes	131.0 MB/s
ChaCha20	ChaCha20	C	128 bytes	194.7 MB/s
ChaCha20	ChaCha20	C	4096 bytes	550.3 MB/s
Legacy /dev/random	SHA-1	C	10 bytes	12.9 MB/s
Legacy /dev/random	ChaCha20	C	16 bytes	29.2 MB/s
Legacy /dev/random	ChaCha20	C	32 bytes	58.6 MB/s
Legacy /dev/random	ChaCha20	C	64 bytes	80.0 MB/s
Legacy /dev/random	ChaCha20	C	128 bytes	118.7 MB/s
Legacy /dev/random	ChaCha20	C	4096 bytes	220.2 MB/s

LRNG - Resources

- Code / Tests / Documentation: <https://github.com/smuellerDD/lrng>
- Testing conducted on
 - Intel x86, AMD, ARM, MIPS, POWER LE / BE, IBM Z, RISC-V
 - Embedded systems and Big Iron
- Backport patches available
 - LTS: 5.10, 5.4, 4.19, 4.14, 4.4
 - 5.8, 4.12, 4.10
- Why is it not upstream?

```
$ cat /proc/lrng_type
DRNG name: drbg_nopr_ctr_aes256
Hash for reading entropy pool: sha512
Hash for operating aux entropy pool: sha512
LRNG security strength in bits: 256
per-CPU interrupt collection size: 1024
number of DRNG instances: 8
Standards compliance: SP800-90B SP800-90C NTG.1
High-resolution timer: true
LRNG minimally seeded: true
LRNG fully seeded: true
Continuous compression: true
```